

✓ Finetuning a pretrained YOLO detection model with Bboxes only:

✓ Imports:

```
1 import os
2 from datetime import datetime
3 import pytz
4 #import tensorflow as tf
```

✓ Some config:

```
1 setupnumber = '_alle' # ONLY for saving the model! <-- TODO: cp noch pythonisieren? :-)
2
3 modeltype = "yolov8n.pt"
4 # Other models:
5 # "yolov8m-seg.pt"
6 # "yolov8s.pt"
7
8 # Train on bounding boxes or segmentations?
9 #labeltype = "boxes"
10 labeltype = 'segs'
11 if labeltype == 'segs':
12     modeltype = modeltype.replace(".pt", "-seg.pt")
13
14 # Bigger batchsizes for smaller models:
15 batchsize = 10
16 if modeltype == "yolov8s.pt":
17     batchsize = 40
18 if modeltype == "yolov8n.pt":
19     batchsize = 120
20 if modeltype == "yolov8n-seg.pt":
21     batchsize = 100
22
23 imagesize = 640 # 1280
24 if imagesize == 1280:
25     batchsize = batchsize // 4
26
27 N_epochs = 300
28
29 localtime = datetime.now(pytz.timezone('Europe/Berlin'))
30 starttime = localtime.strftime("%Y-%m-%d_%H%M")
31 savestring = "Training date = "+starttime
32 savestring += "\nSetup = "+setupnumber
33 savestring += "\nModel = "+modeltype
34 savestring += "\nLabeltype = "+labeltype
35 savestring += "\nBatchsize = "+ str(batchsize)
36 savestring += "\nImagesize = "+ str(imagesize)
37 savestring += "\nEpochs = " + str(N_epochs)
38 print(savestring)
39
```

```
🔗 Training date = 2024-09-28_2146
Setup = _alle
Model = yolov8n-seg.pt
Labeltype = segs
Batchsize = 120
Imagesize = 640
Epochs = 300
```

✓ Fetch the data:

```
1 # Mount Google Drive and copy data:
2 if 'data_loaded' not in globals():
3     from google.colab import drive
4     drive.mount('/content/drive')
5     !mkdir -v /content/Data
6
7     !cp -r /content/drive/MyDrive/setup_015/* /content/Data/ #| pv
8     print("setup015 copied")
9     !cp -r /content/drive/MyDrive/setup_002/* /content/Data/
10    print("setup002 copied")
11    !cp -r /content/drive/MyDrive/setup_003/* /content/Data/
12    print("setup003 copied")
```

```

13 !cp -r /content/drive/MyDrive/setup_006/* /content/Data/
14 print("setup006 copied")
15 !cp -v /content/drive/MyDrive/setup-002-standalone.yaml /content/Data/
16
17 print("Data download finished at", datetime.now(pytz.timezone('Europe/Berlin')).strftime("%H:%M"))
18
↳ Mounted at /content/drive
mkdir: created directory '/content/Data'
setup015 copied
setup002 copied
setup003 copied
setup006 copied
'/content/drive/MyDrive/setup-002-standalone.yaml' -> '/content/Data/setup-002-standalone.yaml'
Data download finished at 22:01

```

✓ Fetch helper scripts/functions and restructure the data:

```

1 if 'data_loaded' not in globals():
2     !cp -v /content/drive/MyDrive/prepare_yolo_data_folders.py /content/
3     # Download helper script from Github: -> doesn't work.
4     #!wget https://raw.githubusercontent.com/<username>/<reponame>/<branch>/prepare_yolo_data_folders.py
5     from prepare_yolo_data_folders import prepare_yolo_data_folders
6
7     # Colab magic to make "Data" the current directory:
8     %cd /content/Data
9
10    #Do the train-validation split and move the image and label files to YOLO folder structure:
11    data_loaded = prepare_yolo_data_folders(labeltype) # TODO: Quiet! :-)
12    print('data_loaded is', data_loaded)

```

↳ '/content/drive/MyDrive/prepare_yolo_data_folders.py' -> '/content/prepare_yolo_data_folders.py'

```

/content/Data
Preparing files and folders for training with segs labels.
Current folder is: /content/Data
Found 354 labels.
Found 354 labelled images.
Created train/images and train/labels.
Created val/images and val/labels.
Created test/images and test/labels.
Moved 247 images and labels to folders.
Moved 71 images and labels to folders.
Moved 36 images and labels to folders.
Data handling finished.
data_loaded is True

```

```

1 # Count the files:
2 num_train = len(os.listdir('/content/Data/train/labels'))
3 num_val = len(os.listdir('/content/Data/val/labels'))
4 num_test = len(os.listdir('/content/Data/test/labels'))
5 print(f"Number of training images: {num_train}")
6 print(f"Number of validation images: {num_val}")
7 print(f"Number of test images: {num_test}")
8
9 savestring += "\nNumber of training images: " + str(num_train)
10 savestring += "\nNumber of validation images: " + str(num_val)
11 savestring += "\nNumber of test images: " + str(num_test)

```

↳ Number of training images: 247
Number of validation images: 71
Number of test images: 36

```

1 # Create model directory on Drive:
2 modeldir = '/content/drive/MyDrive/CucumberModels/'+starttime+"_"+labeltype
3 if not os.path.exists(modeldir):
4     os.makedirs(modeldir)
5
6 savestring += "\nModel directory: " + modeldir

```

✓ Write settings file:

```


1 # Output settings file:
2 settingsfile = modeldir + "_settings.txt"
3 with open(settingsfile, "w") as file:
4     file.write(savestring)
5 print("Model settings written to", settingsfile)

```

↳ Model settings written to /content/drive/MyDrive/CucumberModels/2024-09-28_2146_settings.txt

✓ Install ultralytics and import YOLO:

```
1 %pip -q install ultralytics
2 from ultralytics import YOLO
```

875.1/875.1 kB 49.2 MB/s eta 0:00:00
 Creating new Ultralytics Settings v0.0.6 file 
 View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
 Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see <https://docs.ultralytics.com>

Okay, so the images and labels are there? Next: Import YOLO model for training! Note: Check Bboxes vs. segmentations!

✓ Load tensorboard:

```
1 %load_ext tensorboard
```

✓ Import YOLO model for training

```
1 # Load a pretrained model:
2 model = YOLO(modeltype)
```

Downloading <https://github.com/ultralytics/assets/releases/download/v8.2.0/yolov8n-seg.pt> to 'yolov8n-seg.pt'...
 100%|██████████| 6.74M/6.74M [00:00<00:00, 379MB/s]

✓ YAML file for folder and class definition: Everything ok?






```
1 #!cat /content/misc/setup-002-standalone.yaml
2 yamlfile = "/content/Data/setup-002-standalone.yaml"
```

✓ Do the training!

```
1 # If GPU-RAM is full before training:
2
3 # import torch
4 # torch.cuda.empty_cache()
5
```

```
1 # Train the model
2 trainresults = model.train(data=yamlfile, epochs=N_epochs, imgsz=imagesize, batch=batchsize)
```



Epoch 56/300	GPU_mem 20.9G Class	box_loss 0.7551 Images	seg_loss 1.026 Instances	cls_loss 0.515 Box(P	dfl_loss 0.9179 R	Instances 86 mAP50	Size 640: mAP50-95)	100%  3/3 [00:09<ma
Epoch 57/300	GPU_mem 22.8G Class	box_loss 0.7333 Images	seg_loss 0.9395 Instances	cls_loss 0.5089 Box(P	dfl_loss 0.8963 R	Instances 114 mAP50	Size 640: mAP50-95)	100%  3/3 [00:09<ma
Epoch 58/300	GPU_mem 22.6G Class	box_loss 0.7973 Images	seg_loss 1.003 Instances	cls_loss 0.5384 Box(P	dfl_loss 0.9349 R	Instances 84 mAP50	Size 640: mAP50-95)	100%  3/3 [00:04<ma
Epoch 59/300	GPU_mem 20.8G Class	box_loss 0.7356 Images	seg_loss 0.9242 Instances	cls_loss 0.4935 Box(P	dfl_loss 0.9118 R	Instances 93 mAP50	Size 640: mAP50-95)	100%  3/3 [00:04<ma
Epoch 60/300	GPU_mem 21.8G Class	box_loss 0.743 Images	seg_loss 0.9185 Instances	cls_loss 0.509 Box(P	dfl_loss 0.9054 R	Instances 128 mAP50	Size 640: mAP50-95)	100%  3/3 [00:05<ma

Save the model directly to Drive:

```
1 model.save(modeldir+"/"+starttime+"_"+labeltype+".pt")
```

wandb?

... and managing the parameter "matrix"

> Tensorboard:

[] ↳ 1 Zelle ausgeblendet

✓ Predict the cucumbers on test images:

```
1 # Examples:
2 # #predictresults = model(["image1.jpg", "image2.jpg"]) # return a list of Results objects
3 # # Run inference on 'bus.jpg' with arguments
4 # model.predict("bus.jpg", save=True, imgsiz=320, conf=0.5)
5
6
7 # /content/Data/test/images
8
9 # First, I guess we should _dynamically_?! define a list of images to predict on:
10 #imagelist = []
11 imagelist = ['test/images/'+file for file in os.listdir("/content/Data/test/images")]
12 print(f"Using {len(imagelist)} images for prediction: ")
13 print(imagelist)
14
15 # Then, we can predict on the images:
16 results = model.predict(imagelist,
17                          imgsiz=imagesize,
18                          conf=0.6,
19                          #show=True,
20                          save=True)
21
22 for index, result in enumerate(results):
23     result.save(modeldir+"/"+starttime+"_prediction_"+str(index)+".jpg")
24
25 # # Process result:
26 # # for result in results:
27 # boxes = result.bboxes # Boxes object for bounding box outputs
28 # masks = result.masks # Masks object for segmentation masks outputs
29 # keypoints = result.keypoints # Keypoints object for pose outputs
30 # probs = result.probs # Probs object for classification outputs
31 # obb = result.obb # Oriented boxes object for OBB outputs
32 # result.show() # display to screen

↳ Using 36 images for prediction:
['test/images/dsc01853.png', 'test/images/dsc01868.png', 'test/images/20220614_131249.png', 'test/images/dsc02013.png',
0: 1280x1280 6 cucumbers, 12.0ms
```

```

1: 1280x1280 2 cucumbers, 12.0ms
2: 1280x1280 28 cucumbers, 12.0ms
3: 1280x1280 8 cucumbers, 12.0ms
4: 1280x1280 9 cucumbers, 12.0ms
5: 1280x1280 21 cucumbers, 12.0ms
6: 1280x1280 8 cucumbers, 12.0ms
7: 1280x1280 3 cucumbers, 12.0ms
8: 1280x1280 3 cucumbers, 12.0ms
9: 1280x1280 8 cucumbers, 12.0ms
10: 1280x1280 15 cucumbers, 12.0ms
11: 1280x1280 11 cucumbers, 12.0ms
12: 1280x1280 1 cucumber, 12.0ms
13: 1280x1280 7 cucumbers, 12.0ms
14: 1280x1280 19 cucumbers, 12.0ms
15: 1280x1280 3 cucumbers, 12.0ms
16: 1280x1280 19 cucumbers, 12.0ms
17: 1280x1280 6 cucumbers, 12.0ms
18: 1280x1280 6 cucumbers, 12.0ms
19: 1280x1280 7 cucumbers, 12.0ms
20: 1280x1280 15 cucumbers, 12.0ms
21: 1280x1280 26 cucumbers, 12.0ms
22: 1280x1280 16 cucumbers, 12.0ms
23: 1280x1280 10 cucumbers, 12.0ms
24: 1280x1280 23 cucumbers, 12.0ms
25: 1280x1280 7 cucumbers, 12.0ms
26: 1280x1280 6 cucumbers, 12.0ms
27: 1280x1280 11 cucumbers, 12.0ms
28: 1280x1280 8 cucumbers, 12.0ms
29: 1280x1280 2 cucumbers, 12.0ms
30: 1280x1280 2 cucumbers, 12.0ms
31: 1280x1280 7 cucumbers, 12.0ms
32: 1280x1280 11 cucumbers, 12.0ms
33: 1280x1280 2 cucumbers, 12.0ms
34: 1280x1280 14 cucumbers, 12.0ms
35: 1280x1280 15 cucumbers, 12.0ms
Speed: 166.1ms preprocess, 12.0ms inference, 12.3ms postprocess per image at shape (1, 3, 1280, 1280)
Results saved to runs/segment/train2
Where are the resulting pictures stored? :-)
And how can I rescale them for/after prediction?

```

▼ Show resulting pictures:

Perhaps replace IPython display by PIL stuff?

```

1 #from IPython.display import Image
2 #Image(filename='/content/runs/detect/train53/20220614_134218_Gurkenfoto_2.jpg', width=400)
3

1 #Image(filename='/content/runs/detect/train2/20220614_134218_Gurkenfoto_rescaled_640.jpg', width=400)
2

1 #Image(filename='/content/runs/detect/train53/dsc01843.png', width=400)

```